

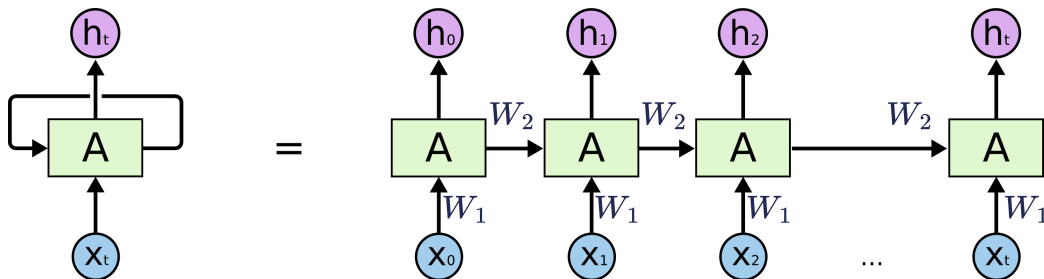
RÉSEAUX DE NEURONES RÉCURRENTS

Vincent Guigue,
inspiré des supports de Nicolas Baskiotis & Benjamin Piwowski





Recurrent Neural Network



General RNN:


$$h_t = \tanh(x_t W_1 + h_{t-1} W_2 + b) \text{ or } \tanh((x_t \oplus h_{t-1})W + b)$$

- **Few parameters**
- Ability to deal with **variable lengths**
- Capture the **dynamics** of the sequence



Cell details

Classical RNN :



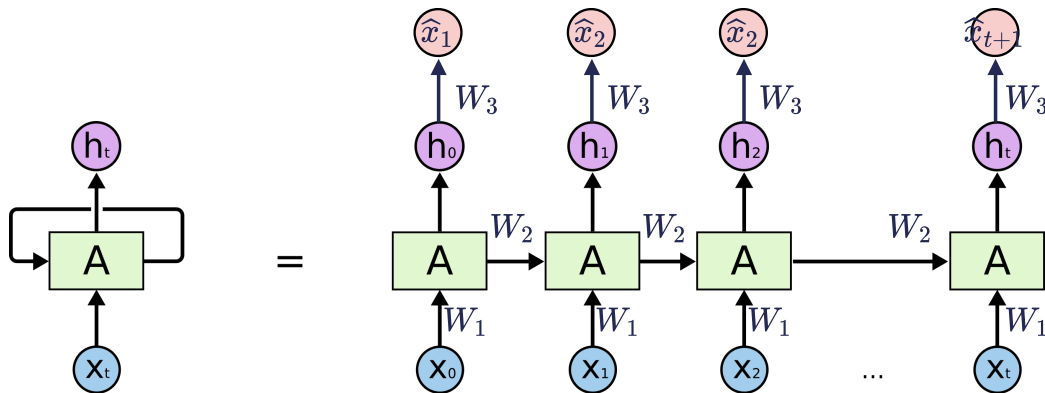
- Latent state $h_t \in \mathbb{R}^d$
- $h_t = f_W(h_{t-1}, x_t) = \tanh((x_t \oplus h_{t-1})W + b)$
- input $x_t \in \mathbb{R}^n$
- $W \in \mathbb{R}^{n \times d}$, $b \in \mathbb{R}$
- Initial state: h_1
- Sequence $[x_1, \dots, x_T] \Rightarrow [h_1, \dots, h_T]$
- May be computed left \rightarrow right (h_t) ... Or right \rightarrow left (h_t^R)



RNN... For which purpose?


- Predicting the next step (from h_T)
- Generation = prediction (in loop)
- Classifying a sequence (from h_T)
- Detecting an event (at every h_t)

Bi-RNN: use h_t & h_t^R
(h_T & h_1^R = caract. whole sequence)





Learn to generate



■ Teacher forcing:

$$\hat{x}_{t+1} = g(h_t), \quad h_{t+1} = f_W(h_t, x_{t+1})$$

■ Free Generation

$$\hat{x}_{t+1} = g(h_t), \quad h_{t+1} = f_W(h_t, \hat{x}_{t+1})$$

In practice:

- Start with teacher forcing (more stable) \Rightarrow then switch to free generation
- Sequential learning also = more complicated solution

\Rightarrow **Reinforcement Learning**



RNN In practice

```
1 # Define the model
2 model = torch.nn.RNN(input_size=100, hidden_size=50,
3                       num_layers=1,
4                       nonlinearity='tanh', bias=True)
5
6 # data (x_1,...,x_T)
7 # Seq. length T=23 & batch of size 10
8 # Input x_t \in R^n dimension 100
9 seq_x = torch.rand(23, 10, 100)
10
11 # seq_h is a tensor (23, 10, 50)
12 # For classification purpose, we use last_h
13 seq_h, last_h = model(seq_x)
```

Convention

Tensor of size **time x sample x representation**

SPECIFICITIES & DIFFICULTIES



Problems with RNN

Simple use case, linear + no entries

- $h_{t+1} = Wh_t$
- t step \Rightarrow computing W^t
- Given the decomposition: $W = U\Sigma U^T$, $\Sigma = \text{diag}$ matrix of eigen values

Gradient computation leads to:

$$\frac{\partial E}{\partial h_t} = \frac{\partial E}{\partial h_N} (W^T)^{N-t} = \Delta U \begin{pmatrix} \sigma_1^{N-t} & & \\ & \ddots & \\ & & \sigma_n^{N-t} \end{pmatrix}$$

- if $\sigma_i > 1$ exponential growth, exploding gradient
- if $\sigma_i < 1$ gradient vanishing



Exploding gradient

- Gradient values > 1 at every iteration
- Weak performances
- *NaN* everywhere quickly !

⇒ **Gradient clipping**

$$\nabla_w E = \begin{cases} \frac{\nabla_w E}{\|\nabla_w E\|} & \text{if: } \|\nabla_w E\| > 1 \\ \nabla_w E & \text{else} \end{cases}$$

```
1 # use before 'optimizer.step()'
2 nn.utils.clip_grad_norm_(
3     model.parameters(),
4     max_norm=1.,
5     norm_type=2
6 )
```




Vanishing gradient \Rightarrow Residual architecture (e.g. ResNet)

$$f(x) = x + f_0(x)$$
$$\frac{\partial f(x)}{\partial x_j} f(x) = \mathbf{1} + \frac{\partial f_0(x)}{\partial x_j}$$

In the back propagation process:

$$\frac{\partial E_f(x)}{\partial x_j} f(x) = \frac{\partial E_y}{\partial y_j} + \sum_i \frac{\partial E_y}{\partial y_i} \frac{\partial y_i}{\partial x_j}$$






Vanishing gradient \Rightarrow gated unit (LSTM)

The phenomenon has been understood & (partially) overcome:

Neurons **learn** what should be **kept in memory** and what should be **forgotten**



Gated architecture



S. Hochreiter, J. Schmidhuber, Neural computation 1997
Long short-term memory

Chris Olah's blog <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>




Vanishing gradient \Rightarrow gated unit (LSTM)

The phenomenon has been understood & (partially) overcome:

Neurons **learn** what should be **kept in memory** and what should be **forgotten**

- Memory c_t (modified by simple operators)
- External representation h_t extracted from c_t
- Gates $y = \sigma(p) \otimes x$,

\otimes : term by term product & σ : sigmoid






Vanishing gradient \Rightarrow gated unit (LSTM)

The phenomenon has been understood & (partially) overcome:

Neurons **learn** what should be **kept in memory** and what should be **forgotten**

- Memory c_t (modified by simple operators)
- External representation h_t extracted from c_t
- Gates $y = \sigma(p) \otimes x$,
 \otimes : term by term product & σ : sigmoid



Gate 1 (forgetting):

$$g_t^{(f)} = \sigma[W_f(x_t \oplus h_{t-1}) + b_f]$$

$$c_t^{(f)} = g_t^{(f)} \otimes c_{t-1}$$




Vanishing gradient \Rightarrow gated unit (LSTM)

The phenomenon has been understood & (partially) overcome:

Neurons **learn** what should be **kept in memory** and what should be **forgotten**

- Memory c_t (modified by simple operators)
- External representation h_t extracted from c_t
- Gates $y = \sigma(p) \otimes x$,
 \otimes : term by term product & σ : sigmoid



Gate 2 (memory update):

$$g_t(i) = \sigma[W_{i,g}(x_t \oplus h_{t-1}) + b_{i,g}]$$

$$v_t(i) = \tanh[W_{i,v}(x_t \oplus h_{t-1}) + b_{i,v}]$$

$$c_t = c_t^{(f)} + g_t^{(i)} \otimes v_t^{(i)}$$




Vanishing gradient \Rightarrow gated unit (LSTM)

The phenomenon has been understood & (partially) overcome:

Neurons **learn** what should be **kept in memory** and what should be **forgotten**

- Memory c_t (modified by simple operators)
- External representation h_t extracted from c_t
- Gates $y = \sigma(p) \otimes x$,

\otimes : term by term product & σ : sigmoid



Gate 3 (output):

$$g_t(o) = \sigma[W_{o,g}(x_t \oplus h_{t-1}) + b_{o,g}]$$

$$v_t(o) = \sigma[W_{o,v}c_t + b_{o,v}]$$

$$h_t = g_t^{(o)} \otimes v_t^{(o)}$$



Conclusion

- **LSTM** and variation (**GRU**) are widely used
- biRNN (biLSTM / BiGRU) offer often faster convergence
- Gradient is always an issue: **watch it!**

```
1 # Calcul de || \nabla_W E ||^2
2 with torch.no_grad():
3     grad_sq12 = sum((p.grad ** 2).sum() for p in l.parameters())
```


- **Dropout** is useful to fight against overfitting
- Layers can be stacked (!) \Rightarrow num_layers

```
1 # Module definition
2 model = torch.nn.LSTM(input_size=100, hidden_size=50,
3                       batch_first=False, num_layers=1,
4                       dropout=.3)
```

DATA GENERATION



Basic data generation




Solution 1: word prediction

- $W_3 \approx$ vocabulary classification
- Several variation: argmax / sampling...
⇒ reduce vocabulary
 - at the letter scale (dim ≈ 60)
 - byte pair encoding \Rightarrow most freq. ngrams of letters
- Solution 1 works better (enable beam search)



Basic data generation



- $W_3 \approx$ vocabulary classification
- Several variation: argmax / sampling...
- ⇒ reduce vocabulary
 - at the letter scale (dim ≈ 60)
 - byte pair encoding \Rightarrow most freq. ngrams of letters
- Solution 1 works better (enable beam search)

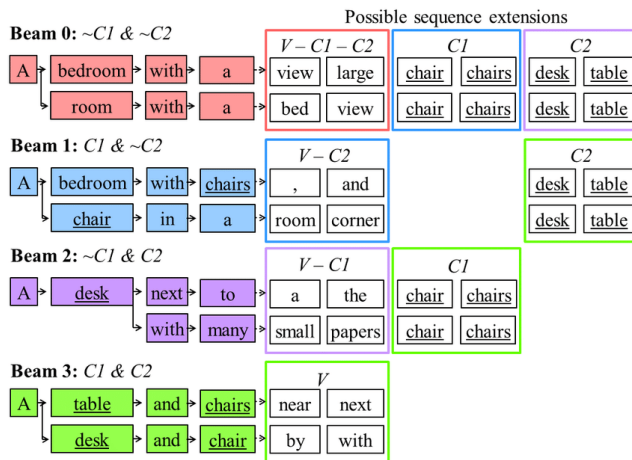
Solution 2: embedding prediction



Beam search

Very difficult to find the best word at time $t...$

⇒ explore several possibilities, keep the most likely !






Evaluation: BLEU/ROUGE (translation & summary)

How to evaluate text generation?

Very difficult...

still an open issue !



- f1 computations
- on N-grams, on longest common subsequence...




Evaluation: BLEU/ROUGE (translation & summary)

How to evaluate text generation?

Very difficult...

still an open issue !



$$2 * \frac{0.43 * 1.0}{0.43 + 1.0} = 0.6$$

60% f1 score

- f1 computations
- on N-grams, on longest common subsequence...



Conclusion

- Evaluation in the latent-space (cf BERT-score / CLS)
- Content evaluation (entity, figures...) \Rightarrow hallucination problems